

Ollama Guide

Author: CarlMart

Date: Monday, July 27, 2026

Abstract: A Step-by-step User Guide for Installing and running Ollama on terminal window

Work in Progress



TOC

- [Chapters](#)
 - [Docker](#)
 - [ollama on macOS](#)
 - [Installation](#)
 - [Models](#)
 - [Token Counts](#)
 - [Edit: Work in progress](#)

Do ollama models grow and learn as they continue to interact with them?

Ollama is an AI model based on the OpenAI's Qwen language model. It works similarly to

how a human model could interact with Ollama: it starts by asking for basic information, such as what type of question you want answered or provide some context about.

However, when interacting with Ollama over time, Ollama will continue to learn and improve based on the data it receives. It can adapt its language patterns and responses, improving its ability to understand natural language queries and generate more relevant and accurate answers in future interactions.

Ollama uses algorithms like reinforcement learning to improve itself over time. It learns from past interactions with you, using the feedback to make its next interaction better than any previous one.

Docker

Docker is an open-source platform that allows developers to build, ship, and run applications in standardized, lightweight units called containers.

Installing and running it makes things somewhat easier to maintain but overall installing ollama without docker is simpler for daily use.

Both methods of deploying will be shown and choose the method that works best for you.

Choosing a Model

A basic Model to run depends on the amount of Ram you currently run on your PC.

The Table below will help you decide on which model to choose.

If you run NVIDIA card , you can utilize your GPU to run faster queries.

Max RAM	Model	Notes
8GB-4GB	gemma3n	-----
8GB-4GB	llama3.2:3b	-----
16GB-6-8GB	llama3.1:8b	-----
16GB-6-8GB	qwen2.5-coder:7b	for Programmers
16GB NVIDIA	Llama 3.1 8B	-----
16GB NVIDIA	Qwen 2.5 7B	-----

Smallest ollama models

Model	Usage	Notes
qwen2.5:0.5b	txt,chat,multilingual	398MB
qwen2.5-coder:0.5b	lightweight coder	-----
llama3.2:1b	slightly large	1.3GB

"Avoid models larger than 8B to prevent severe slowdowns.

Qwen 2.5 7B at Q4_K_M quantization uses roughly 5.2 GB of RAM,

leaves 9 GB for your OS and apps, runs at 8-12 tokens/sec on CPU only

35-50 tokens/sec on a midrange GPU, and benchmarks within 4% of Llama 3.1 8B

on most tasks while being noticeably faster."

- [ref. reddit:LocalLLaMa](#)

Model specs

- [smallest model: smollm2](#)
- [coder model: qwen2.5-coder](#)

Installation

Installs can be direct app or in a terminal window.

Its a quick install since its only the wrapper app.

- ollama.com/download
- [Windows and Mac install](#)

```
curl -fsSL https://ollama.com/install.sh | sh # mac /Linux  
irm https://ollama.com/install.ps1 | iex | # Windows
```

macOS - default model directory

```
~/ollama/models  
~/ollama/models/manifests/registry.ollama.ai/library/PERSONAL
```

Location contains two main subdirectories:

```
blobs : hold actual model weights as binary files named by SHA256  
hashes  
manifests : contain JSON metadata linking model names to their  
corresponding blobs
```

Storage location:

```
set the OLLAMA_MODELS environment variable to new path  
export OLLAMA_MODELS=/mnt/NFS/ollama/models/  
use a symbolic link.  
cd ~/.ollama/ ; mv models /mnt/NFS/ollama/  
ln -s /mnt/NFS/ollama/models/ ~/.ollama/models
```

CLI Administration

All you need is one model . The commands are shown below

The models shown below are larger GB size models so will take some time depending

on your download speed.

The commands to run are shown below. You only pull one model , but you can download as many as you want. You will only run one instance.

```
ollama serve
ollama pull qwen3:8b
ollama pull Llama 3.2:3B
ollama pull qwen3:8b
ollama pull qwen2.5-coder          - Download : smaller
version
ollama pull qwen2.5-coder:7b       - Download : good for 16GB
ollama pull qwen2.5:7b-instruct-q4_K_M -
ollama run qwen2.5:7b-instruct-q4_K_M -
ollama run <model>
ollama run <model> | tee chat_history.txt
ollama run qwen3.5:0.8b "hello world" --think=false
-- false means just the answer , no rambling
```

A few usefull commands

```
ollama run qwen3.5:0.8b "hello world" --think=false
-- for terminal formatted markdown

ollama list          - view downloaded models
ollama rm <model>    - delete model
ollama ps            - view whats running
ollama --version      - self explanatory
ollama stop <model>  - to completely stop model
```

Interactive commands

```
> /?      - help
> /bye    - to quit the interface , keeps running for 5 min unless you
stop it
  /save <filedir>
  /load <filedir>
```

When saving chat, files on macOS are stored at

```
~/.ollama/models/manifests/registry.ollama.ai/library/<filedir>
```

Ask-ollama

Installation for console app

cp ~/.cargo/bin/ask ~/binapps/ or create a shell script ask

```
cargo install ask-ollama
```

default model is mistral so make sure you specify model

```
~/.cargo/bin/ask "what is capitol of france"  
ask --model=qwen2.5:7b-instruct-q4_K_M "question"  
ask --version  
ask --help
```

Ask-ollama-JSON

Custom console app - used for import/export JSON chats.

Be sure to install '**md2term**' to view output in nicely formatted markdown output to terminal

```
pip install md2term  
ask-ollama-json.py
```

Sample session with ask-ollama-json.py

Ollama Interactive TUI Chat Manager

Detected Local Ollama Models

Index	Model Name
1	smollm2:latest

Enter model name or number index (1):

Do you want to load a previous chat JSON file? [y/n]: n

--- Chat active with model: smollm2:latest ---

Commands: Type exit to save/quit | Type /clear to wipe screen.

You: what is the area of a circle?

Assistant: The area of a circle can be calculated using the formula:

$$\text{Area} = \pi * r^2$$

Where " π " (pi) is approximately 3.14159 and "r" is the radius of the circle...

You: exit

Assistant: Would you like to save this conversation? [y/n]: y

Enter filename to save as (chat_history.json): math.json

Chat history safely exported to 'math.json'.

open-webui

Used for uploading pdf + other doc files to summarize

Use conda activate ollama. md2term is to display markdown to a terminal window nicely formatted

Install in a python environment like conda or venv

```
pip install open-webui md2term rich requests
```

Python environment using conda

First list all environments you have created, then activate or deactivate
python=3.11 was installed on ollama environment (2026)

```
conda info --envs          * lists created environments
conda env export -n ollama --from-history * lists dependencies
conda activate ollama      * activates it
pip install openai         *
```

Starting open-webui

There are two methods to start open WebUI

1. Desktop Application : open /Applications/Open\ WebUI.app
2. Browser : More complicated , but uses less resources

```
screen -S ollama
conda activate ollama
open-webui serve          -- to start baremetal
or
docker run -d -p 3000:8080 -v open-webui:/app/backend/data --name
open-webui ghcr.io/open-webui/open-webui:main
ctrl-A , ctrl-D to leave screen
pkill -f "open-webui serve" -- to stop server
```

Once running access UI :

```
http://localhost:8080      - for open-webui serve
http://localhost:3000
```

Ollama terminal interface

md2term displays markdown text to formatted terminal

```
ollama run qwen2.5:7b-instruct-q4_K_M "hello world"
ollama run qwen2.5:7b-instruct-q4_K_M "hello world" --think=false
ollama run qwen2.5:7b-instruct-q4_K_M "hello world" --think=false |
md2term
```

Web gui and Desktop app

- [native desktop app 4 Windows/Linux/Macos](#)
- [native desktop app Docs](#)
- [inkeybit blog openwebui-ollama guide](#)
- [internet search ddgs free](#)

WebUIs: Interfaces like Open WebUI automatically save chat histories to a database, allowing you to reload them later. Export is in JSON format

When reloaded, the interface feeds the old messages back to the model to reconstruct the context.

Install Docker first if not installed for web gui

- [ollama and webui](#)
- [localhost:3000](#)

Docker scripts v1

```
docker run -d \  
  -p 3000:8080 \  
  --add-host=host.docker.internal:host-gateway \  
  -v open-webui:/app/backend/data \  
  --name open-webui \  
  --restart always \  
  ghcr.io/open-webui/open-webui:main
```

Docker scripts v2

```
docker run -d --network=host \  
  -e OLLAMA_BASE_URL=http://127.0.0.1:11434 \  
  -v open-webui:/app/backend/data \  
  -p 3000:8080 \  
  --name open-webui \  
  --restart always \  
  ghcr.io/open-webui/open-webui:main
```

open-webui to ollama connections

- [connection fix](#)

```
# Step 1: Is Ollama running at all?
curl http://localhost:11434
# Should print "Ollama is running"

# Step 2: Does Ollama have models?
curl http://localhost:11434/api/tags
# Should print JSON with your model list

# Step 3: Can Open WebUI's container reach Ollama?
docker exec open-webui curl http://host.docker.internal:11434
# should print
100      17   100      17      0      0   3839      0  --:--:--  --:--:--  --:--
:--  4250Ollama is running
# Or if using --network=host:
docker exec open-webui curl http://127.0.0.1:11434
```

More apps

- [Openclaw.ai assistant](#)

OpenClaw is a self-hosted gateway that connects your favorite chat apps — Discord, Google Chat, iMessage, Matrix, Microsoft Teams, Signal, Slack, Telegram, WhatsApp, Zalo, and more via channel plugins — to AI coding agents

```
npm install openclaw
```

Free CLI Tools

tgpt - supports engines like Phind without requiring an API key or login.

tgpt and similar open-source projects provide free access to AI capabilities directly in the terminal.

Savvy is mentioned as a free CLI tool for individual developers that does not require API keys.

chatGPT-shell-cli require an OpenAI API key,

- [free alts to copilot cli](#)
- [chatgpt command line alt](#)
- [chatgpt shell cli](#)

llama.cpp llama-server install

Open-WebUI uses port 8080

so for llama.cpp server port use 8081 :

http://127.0.0.1:8081

Note: Ensure the port 8081 does not conflict with other services.

```
/usr/local/bin/llama-server -m <path>/<model>.gguf --port 8081
```

llama-server run

Running llama-server with model then connecting to open WebUI

Point Open WebUI to llama.cpp server

```
Open WebUI -> Admin -> Settings -> Connections  
add: URL: http://localhost:8081/v1
```

[ref OpenWebUI to llama.cpp](#)

llama-server start

With or Without API Key : --api-key "123" - some demo key
or docker

```
llama-cli -hf bartowski/Meta-Llama-3.1-8B-Instruct-GGUF:Q4_K_M
llama-server -m ~/.ollama/models/gguf/model.gguf --port 8081
llama-server -m ~/.ollama/models/gguf/smollm2/smollm2-1.7B-Q8_0.gguf -
-port 8081 --api-key "123"
docker run -d --network host -v open-webui:/app/backend/data --name
open-webui ghcr.io/open-webui/open-webui:main
```

Test Connection with curl

```
curl http://127.0.0.1:8081/v1/models -H "Authorization: Bearer sk-
llama-local-abc123"
  or without key
curl http://127.0.0.1:8081/v1/models

  output should be  JSON format
  ...smollm2-1.7B-
Q8_0.gguf", "model": "modified_at": "", "size": "model", "description"....
```

Actual test case and output

```
llama-server -m  ~/.ollama/models/gguf/smollm2/smollm2-1.7B-Q8_0.gguf
--port 8081

curl -s http://127.0.0.1:8081/v1/chat/completions -d \
  '{"model": "mistral", "messages":
[{"role": "user", "content": "Hello"}]}' | \
  jq '.choices[0].message.content'

  > output: "Hello! How can I assist you today?"

http://localhost:3000/?temporary-chat=true  // open WebGUI interface
```

Conversions and Downloads of gguf files

Download file manually from Hugging Face repositories

eg. bartowski, TheBloke, or ggml-org then point llama.cpp -m flag:

-hf = hugging face

- [hugging face gguf-llamacpp](#)

```
llama-cli -m /path/to/your_model.gguf
llama-server -hf bartowski/Llama-3.2-3B-Instruct-GGUF:Q8_0
```

huggingface

```
curl http://localhost:8080/v1/chat/completions \
-H "Content-Type: application/json" \
-H "Authorization: Bearer no-key" \
-d '{
  "messages": [
    {
      "role": "system",
      "content": "You are an AI assistant. Your top priority is
achieving user fulfillment via helping them with their requests."
    },
    {
      "role": "user",
      "content": "Write a limerick about Python exceptions"
    }
  ]
}'
```

Converting ollama to llama.cpp format files

OllamaToGGUF.py - script to convert from Ollama's split format to a single GGUF (GPT General Unified Format) file.

reads model manifests and combines model layers stored in blob files to a GGUF file.

- [github OllamaToGGUF](#)

Simple script to convert blob files to GGUF. Nothing to compile. Just run it.

I renamed it ollama2gguf.py and copied it to local

```
git clone https://github.com/mattjamo/OllamaToGGUF.git
OllamaToGGUF/OllamaToGGUF.py

mkdir -p ~/.ollama/models/gguf/
cd ~/.ollama/models/gguf/
```

You can also just use ollama to show you the exact file:

```
ollama show qwen2.5-coder:7b --modelfile | grep -v \# | grep FROM
FROM ~/.ollama/models/blobs/sha256-
60e05f2100071479f596b964f89f510f057ce397ea22f2833a0cfe029bfc2463
# then just point llama.cpp to exact file
llama-cli -m ~/.ollama/models/blobs/sha256-
60e05f2100071479f596b964f89f510f057ce397ea22f2833a0cfe029bfc2463
```

Running llama.cpp with api key and gguf file

```
llama-server -m ~/.ollama/models/gguf/smollm2/smollm2-1.7B-Q8_0.gguf -
-port 8081 --api-key "sk-llama-local-abc123"
```

Running llama.cpp cli with ollama blobs

```
llama-cli -m ~/.ollama/models/blobs/sha256-  
60e05f210064f89f510f057ce397ea22f2833a0cfe029bfc2463
```

Running llama.cpp with openWebUI and docker

```
docker run -d --network host -v open-webui:/app/backend/data --name  
open-webui ghcr.io/open-webui/open-webui:main
```

Daily ollama startup

ollama starts at startup boot

is ollama running?

```
curl http://localhost:11434  
output : ollama is running
```

does ollama have models?

```
curl http://localhost:11434/api/tags  
output: {"models":[{"name":"smollm2:latest" ... json file with  
model list
```

ollama list

NAME	ID	SIZE	MODIFIED
smollm2:latest	cef4a1e09247	1.8 GB	33 hours ago
qwen2.5-coder:7b	dae161e27b0e	4.7 GB	2 days ago
mistral:latest	6577803aa9a0	4.4 GB	3 days ago
qwen2.5:7b-instruct-q4_K_M	845dbda0ea48	4.7 GB	3 days ago

Starting open-webui

Starting with llama-server and browser access

```
screen -S ollama
llama-server -m ~/.ollama/models/gguf/<model>.gguf --port 8081
open-webui serve - starts baremetal
```

start llama.cpp

```
screen -S llamacpp
llama-server -m ~/.ollama/models/gguf/smollm2/smollm2-1.7B-Q8_0.gguf -
-port 8081 --api-key "sk-llama-local-abc123"
```

access open-webui via browser. login

```
http://localhost:8080/ to access open-webui via browser
```

Complete Screen , Env , llama-serve and open-webui

```
conda activate ollama
llama-server -m $MOD --port 8081 &
# URL: http://127.0.0.1:8081
open-webui serve &
# URL: http://localhost:8080/
```

Import json files

First install a few libraries

```
pip install pandas openpyxl
```

Python code

```
import pandas as pd
df = pd.read_json('https://example.com/your-data.json')

or

import pandas as pd
df = pd.read_json('path/to/your/file.json')
```

Ollama Models with Real-Time Web search

SearXNG is a free, open-source metasearch engine that aggregates results from various search providers without storing any of your data.

First step is to deploy searxng

```
docker run -d --name searxng -p 8080:8080 searxng/searxng:latest
```

Then Open Web UI v

```
docker run -d --name open-webui -p 3000:8080 --add-  
host=host.docker.internal:host-gateway ghcr.io/open-webui/open-  
webui:main
```

Configure Web Search:

- Log in to Open WebUI: **http://localhost:3000**
- Navigate to Settings > Admin Settings > Connections.
- Set Ollama base to docker or bare metal
 - URL: **http://host.docker.internal:11434**
 - URL: **http://localhost:11434**
- Navigate to Settings > Admin Settings > Web Search.
- Enable SearXNG
 - set the URL to SearXNG instance:
 - **http://host.docker.internal:8080/search**
- Save settings. You can now toggle "Web Search" in individual chat sessions.

Verbose output

Tutorial's on Voice Virtual Assistant

- [Build Your Own - medium](#)
- [JARVIS YouTube](#)

TinyLlama

Yes you can install and run tinyllama on R Pi

A few YouTube links:

- [Rapsberry Pi 5 - YouTube](#)
- [Local AI assistant - YouTube](#)

Token Counts

When running your local ollama LLM you can use the --verbose to view output

- [Understanding tokens](#)
- [what are tokens?](#)

```
ollama run tinyllama --verbose
```

AI Agents

AI agents are autonomous software systems designed to perceive their environment, reason through problems, and take independent action to achieve specific goals with minimal human oversight.

AI agents can plan multi-step workflows, adapt to changing conditions, and utilize external tools or APIs to execute complex tasks

- [infor.com What are AI Agents](#)
- [Build a local AI agent](#)

There are several open-source AI agent frameworks and sample projects specifically designed to work with Ollama.

small ai agent to test ollama personal

For a lightweight local AI agent to test **Ollama** on personal hardware, the best options depend on whether you want a finished application or a custom build.

Top Ready-Made Agents

- **Open Interpreter:** Widely recommended for **terminal automation** and coding; it is stable and integrates seamlessly with Ollama models like **Qwen2.5-Coder** or **Llama 3**.
- **Goose:** An open-source agent from Block that connects to **MCP tools** and is ideal for

users who want a finished local agent without coding.

- **OpenClaw**: A stable personal assistant that runs locally via Ollama, supports **WhatsApp/Telegram** channels, and uses models like **Qwen 3.5**.

Minimal Custom Code Example

To understand the mechanics, you can build a simple agent in **~60 lines of Python** that calls Ollama's **OpenAI-compatible endpoint** (`http://localhost:11434/v1`), parses **tool calls**, and executes them locally.

```

import json
from openai import OpenAI

# Connect to local Ollama instance
client = OpenAI(base_url="http://localhost:11434/v1",
api_key="ollama")
MODEL = "qwen3:8b" # Or mistral, llama3.1, etc.

def run_agent(goal: str, max_steps: int = 6):
    messages = [
        {"role": "system", "content": "You are a local AI agent. Use
tools when needed."},
        {"role": "user", "content": goal},
    ]

    for _ in range(max_steps):
        resp = client.chat.completions.create(model=MODEL,
messages=messages, tools=[
            {"type": "function", "function": {"name": "echo",
"parameters": {"type": "object", "properties": {"text": {"type":
"string"}}}}}
        ])
        msg = resp.choices[0].message
        messages.append(msg)

        if not msg.tool_calls:
            return msg.content

        # Execute the tool (simplified for testing)
        call = msg.tool_calls[0]
        result = f"Executed: {call.function.name}"
        messages.append({"role": "tool", "tool_call_id": call.id,
"content": result})

    return "Max steps reached."

print(run_agent("Echo 'Hello World'"))

```

Recommended Models for Testing

- **Qwen3 8B**: Best all-round small agent; requires **~6–8 GB VRAM**.
- **Llama 3.1 8B**: Well-documented with **128K context**; requires **~6–8 GB VRAM**.
- **Mistral Small 3.2**: Best for **clean JSON output** and low-latency function calls; requires **~14–16 GB VRAM**.

Popular Open-Source Frameworks

- **LangChain & LangGraph**: Used for building agents with tool-calling and memory. LangChain provides ChatOllama integration
 - Sample Use Case: Creating a RAG (Retrieval-Augmented Generation) bot that reads local PDFs or a voice assistant with short-term memory.
- **CrewAI**: Focuses on multi-agent systems where different agents have specific roles (e.g., "Researcher" and "Writer") and collaborate to solve tasks.
 - Sample Use Case: An automated blog post generator where one agent researches a topic and another writes the content
- **Goose**: An open-source agent designed for local execution. It excels at developer workflows, such as writing code, running terminal commands, and editing files.
 - Sample Use Case: An autonomous coding assistant that refactors code or debugs errors directly in your IDE.
- **Open Interpreter**: Allows LLMs to run code (Python, JavaScript, etc.) locally to complete tasks. It works seamlessly with Ollama for a fully offline experience
 - Sample Use Case: "Analyze this CSV file and plot a graph of sales trends," where the agent writes and executes the Python script itself.
- **TrustGraph**: A framework with built-in agent flows and graph database integration, supporting Ollama out of the box for complex reasoning tasks.

Model links

- [LangChain/LangGraph](#)
- [CrewAI](#)

- [Goose](#)
- [OpenInterpreter](#)
- [TrustGraph](#)
- [Build Your Own - Ollama and CUA](#)
- [Open-Source Stanford's ACE framework - learns by mistakes](#)

Minimal Ram Requirements

Recommended Models for Agents

Not all models work well for agentic tasks (which require strict tool-calling formats). The following are currently ranked highest for reliability with Ollama:

Model	Best For	VRAM Requirement
Qwen3 (8B/30B)	Best all-round agent performance	6–20 GB
Llama 3 Groq Tool Use	Specialized function calling	~5 GB
Qwen2.5-Coder	Coding and software development agents	4–20 GB
Gemma 4	Native function calling and JSON output	16–21 GB

Latest Desktop releases

- [Github Desktop Releases](#)

Legal Analysis

- [Ultimate guide](#)

Key Applications and Capabilities:

- **Clause Extraction and Structuring:** LLMs parse unstructured legal text (PDF/DOCX) to identify and normalize specific clauses (e.g., **Governing Law**, **Liability Cap**) into structured formats like JSON, enabling faster due diligence and data retrieval. 
- **Contract Comparison and Negotiation:** Advanced approaches use **Natural Language Inference (NLI)** to compare executed contracts against standard templates.  By testing for entailment in both directions, LLMs can identify deviations, flag non-standard clauses with high accuracy, and suggest revisions to template documents.
- **Risk Identification:** LLMs evaluate contracts for legal risks such as **uncapped liability**, **unfair termination rights**, or **IP ownership issues**.  Benchmarks like **ContractEval** demonstrate that both proprietary and open-source models can identify clause-level risks effectively. 
- **Performance vs. Human Review:** Studies indicate that advanced LLMs can match or exceed **junior lawyers** and **Legal Process Outsourcers (LPOs)** in accuracy while completing reviews in seconds rather than hours. 

Leading Models for Legal Contracts:

Leading Models for Legal Contracts:

- **Proprietary Models:** **GPT-4o**, **Claude Opus 4.1**, and **Gemini 2.5 Pro** are widely used for their strong reasoning and long-context capabilities. 
- **Open-Source and Specialized Models:** **DeepSeek-R1** excels in complex reasoning for contract review.  Legal-specific baselines like **Contracts-BERT**, **Legal-BERT**, and **CaseLaw-BERT** often outperform general models on legal tasks despite having fewer parameters.  **Evisort** has also developed a proprietary LLM trained on millions of contracts for specialized contract lifecycle management. 

Implementation Considerations:

- **Accuracy and Hallucination:** While LLMs are powerful, they may hallucinate legal nuances. Human-in-the-loop review remains essential for critical decisions.
- **Data Privacy:** Legal contracts contain sensitive data. Organizations often use **private deployments** or **fine-tuned models** to maintain confidentiality, avoiding the use of public APIs for sensitive documents.
- **Structured Outputs:** Using frameworks like **Pydantic** ensures consistent JSON output, which is critical for integrating LLM results into existing contract management systems. 